

EC2 Pricing Models in AWS

There are 5 main pricing models for EC2:

1. On-Demand Instances

- Pay-as-you-go model: no long-term commitment.
- Best for: short-term workloads or unpredictable traffic.
- Highest cost per hour, but very flexible.

2. Reserved Instances (RIs)

- Commitment-based (1 or 3 years).
- Up to 75% cheaper than On-Demand.
- Three types:
 - **Standard RIs** – fixed usage.
 - **Convertible RIs** – change instance type.
 - **Scheduled RIs** – for specific time windows.

3. Spot Instances

- Unused EC2 capacity at up to 90% discount.
- Can be terminated by AWS anytime if demand increases.
- Ideal for fault-tolerant or batch jobs.

4. Savings Plans

- Commit to a consistent amount of usage (\$/hr) for 1 or 3 years.
- More flexible than RIs (can change instance types, OS, regions).
- Two types:
 - Compute Savings Plan
 - EC2 Instance Savings Plan

5. Dedicated Hosts

- Physical servers fully dedicated to your use.
- Helps meet compliance or licensing needs (e.g., BYOL).
- Expensive but offers complete hardware control.

EC2 Instance Types / Families

EC2 instance types are grouped into **families** based on **use case and hardware optimization**:

| Family | Use Case | Examples |
|------------------------------|---|------------------|
| General Purpose | Balanced CPU, memory, networking | t3, t4g, m5, m6i |
| Compute Optimized | High-performance CPU tasks | c5, c6g, c7i |
| Memory Optimized | In-memory databases, large caches | r5, r6g, x2idn |
| Storage Optimized | High storage I/O (e.g., big data, OLAP) | i3, i4i, d3en |
| Accelerated Computing | GPUs, FPGAs (ML, HPC) | p4, g5, inf2, f1 |
| High Memory | SAP HANA, in-memory DBs | u-6tb1, u-12tb1 |
| Mac Instances | Build, test Apple macOS apps | mac1, mac2 |

Example Format:

Instance type m5.large

- Family: m5 (General purpose)
- Size: large (determines vCPU, RAM)